

# Bert's Restaurant Group — Operations Platform

Prepared September 17, 2026 · Response to  
independent review

Assessment of “Bert Inventory App Review” by Jeff — context, validity, and suggestions

---

## 1. What the document is

A cautious, read-only third-party review of the rebuilt inventory and food-costing application, focused on one question: can this support a crew-to-owner count-and-order workflow? The verdict: a good candidate foundation — the item catalog, storage areas, count entry, supplier pack details, invoice history, order suggestions, and AI panel are all confirmed present — but session-based counting, approvals, and notifications are needed before it is called usable. That framing is sound and we accept it.

## 2. Points that are valid — confirmed against the actual code

### 1. Count entry pre-fills old values (his Issue #3).

True. Enter Counts pre-loads current stock into the inputs, and a blank field saves as zero. There is no “counted / not counted / couldn’t count” state, no record of who counted, and no dated session. His fix — explicit states, completion tracked per storage area, blank must not mean zero — is the right design.

### 2. Mobile layout at 390px.

True, and a regression we introduced: the original app hid secondary columns on narrow screens; the rebuild dropped that rule, so the count table overflows (~435px) on phones. Count entry should become item cards or a step-through form on small screens.

### 3. The AI claimed it added ground beef (his Issue #1).

The most important finding in the document, and it is real. The assistant (“Sous”) has no ability to write anything — it can only read live data and reply. Nothing was actually changed. But its wording implied a stock update had happened, which is dangerous: staff could believe inventory changed when it did not. The fix is to constrain the assistant to advise-only language, and to never confirm a change that did not occur. His principle — requested ≠ ordered ≠ confirmed ≠ received — is exactly right.

### 4. No order history or approval chain.

True. The Order Generator computes suggestions and exports CSVs, but nothing is stored as a dated crew request → owner approval → supplier submission → receipt. His model of keeping requests separate from stock is the correct architecture for that module.

### 5. No authentication.

True — already flagged in our own delivery notes. Anyone with the link can reach the Ownership view and edit items at all three restaurants. Role-based access (staff / manager / ownership, per-restaurant) must exist before staff distribution.

### 6. Operations tab is a placeholder.

Accurate, and intentional — scheduling is the next roadmap module, followed by checklists/compliance, training, and SOPs.

### 3. Where context changes his conclusions

Not wrong — simply unverifiable from a read-only preview. For the record:

- **AI data scope.** He could not verify it, but we did: the assistant queries the live database (inventory, costing, prep, waste) for the selected location and passed end-to-end testing citing actual plate costs and par levels. The data scope is real — the write claims are the problem, not the reads.
- **Seed data.** He reviewed 8 sample items at Bert's, not the real 163-item Product Master. His “map the full order guide” point stands; importing the real product list is already a planned task.
- **Duplicate invoice handling.** Implemented (vendor + invoice number + date matching blocks re-imports) — he simply could not test it read-only.
- **“Connection to this Codex conversation.”** That refers to his own tooling workflow, not an app defect. Only relevant if a Codex integration is ever desired.

### 4. The one item we cannot confirm

His Issue #4 — Storage Area options rendering as “()” and vendor/SKU selectors showing “-”. Nothing in the code produces that with the seeded data, and our automated browser testing did not reproduce it. It may be a stale-preview artifact or a browser-specific rendering issue — but he also observed console warnings about option markup, so the correct stance is: investigate, do not dismiss.

### 5. Suggestions, in priority order

#### Quick fixes — small, high value

- AI guardrail: the assistant explicitly cannot modify data — advise-only, and never confirm unsaved changes.
- Blank count  $\neq$  zero; add counted / not-counted / couldn't-count states.
- Restore a mobile-friendly count entry (card layout under ~500px); test on an actual staff phone.
- Investigate the “()” / “-” selector rendering he reported.

#### Structural — his “smallest useful extension,” which we endorse

- Dated count sessions with saved/pending states and immutable submission revisions — corrections create new revisions, never erase the original.
- Completion notifications to ownership with restaurant, person, time, and session link; failed or offline saves must never look complete.
- Order request → approval → supplier submission → receipt stored as separate records; owner comparisons stay stock-vs-stock and order-vs-order.
- Authentication with per-restaurant staff / manager / ownership roles — before any staff distribution.

#### Validation

- Adopt his “first proof on a real phone” test plan as-is once the quick fixes and count sessions land: count real items on one phone (partial case, each-unit, explicit zero, unresolved), verify on a second device, complete a session deliberately, confirm one completion alert, make a correction and verify the original survives, and test interrupted saves without duplicate notifications.

#### Between you and Jeff

- His request for editable project access is your call — if you want him building on this, share the project or export with him directly.

---

Bottom line: the review is fair and largely accurate. Its strongest contributions are the AI write-claim catch, the blank-means-zero count risk, the mobile regression, and the session/approval/notification workflow model — all accepted. Its caveats about AI data scope, duplicate handling, and authentication reflect the limits of a read-only preview rather than defects. Recommended next step: implement the four quick fixes, then build the count-session workflow.